

Spotify Descriptive Statistics

Matthew, Sam, Andrew

Setup

```
library(readxl)
library(dplyr)
library(car)
library(ggplot2)

sp <- read_excel("Data/spotify_stat3120.xls")
```

Descriptive Statistics

Danceability

```
dance_var <- sp$danceability
summary(sp$danceability)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.166	0.511	0.636	0.629	0.754	0.954

```
sd(sp$danceability)
```

```
[1] 0.1594543
```

Energy

```
energy_var <- sp$energy  
summary(sp$energy)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.1090	0.6430	0.8230	0.7782	0.9350	0.9980

```
sd(sp$energy)
```

```
[1] 0.177742
```

Loudness

```
clean_sp <- filter(sp, loudness < 0)  
clean_loud <- clean_sp$loudness  
summary(clean_loud)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-17.912	-8.063	-6.126	-6.471	-4.737	-0.344

```
sd(clean_loud)
```

```
[1] 2.857596
```

```
var(clean_loud)
```

```
[1] 8.165856
```

Tempo

```
tempo_var <- sp$tempo  
summary(sp$tempo)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
81.99	129.98	145.00	147.22	160.18	220.04


```
sd(sp$tempo)
```

```
[1] 23.24385
```

Plots

Danceability

```
layout(mat = matrix(c(1,2),2,1, byrow=TRUE), height = c(1,8))
par(mar=c(0, 3.1, 1.1, 2.1))
boxplot(dance_var, horizontal=TRUE, ylim=c(0,1), xaxt="n",
        col=rgb(0.8,0.8,0,0.5), frame=F)
par(mar=c(4, 5, 1.1, 2.1))
hist(dance_var, breaks=40, col=rgb(0.2,0.8,0.5,0.5), border=F,
     main="", xlab="Danceability", ylab="Frequency", xlim=c(0,1))
```




Energy

```

layout(mat = matrix(c(1,2),2,1, byrow=TRUE), height = c(1,8))
par(mar=c(0, 3.1, 1.1, 2.1))
boxplot(energy_var, horizontal=TRUE, ylim=c(0,1), xaxt="n",
        col=rgb(0.2, 0.4, 0.8, 0.5), frame=F)
par(mar=c(4, 5, 1.1, 2.1))
hist(energy_var, breaks=40, col=rgb(0.2, 0.4, 0.8, 0.8), border=F,
     main="", xlab="Energy", ylab="Frequency", xlim=c(0,1))

```




Loudness

```


layout(mat = matrix(c(1,2),2,1, byrow=TRUE), height = c(1,8))
par(mar=c(0, 3.1, 1.1, 2.1))
boxplot(clean_loud, horizontal=TRUE, ylim=c(-60,0), xaxt="n",
        col=rgb(0.2, 0.7, 0.4, 0.5), frame=F)
par(mar=c(4, 5, 1.1, 2.1))
hist(clean_loud, breaks=40, col=rgb(0.1, 0.5, 0.3, 0.8), border=F,
     main="", xlab="Loudness", ylab="Frequency", xlim=c(-60,0))

```



Tempo

```
layout(mat = matrix(c(1,2),2,1, byrow=TRUE), height = c(1,8))
par(mar=c(0, 3.1, 1.1, 2.1))
boxplot(tempo_var, horizontal=TRUE, ylim=c(50,250), xaxt="n",
        col=rgb(0.6, 0.3, 0.8, 0.5), frame=F)
par(mar=c(4, 5, 1.1, 2.1))
hist(tempo_var, breaks=40, col=rgb(0.7, 0.2, 0.6, 0.8), border=F,
     main="", xlab="Tempo", ylab="Frequency", xlim=c(50,250))
```



Multiple Linear Regression (predict danceability)

```
lm_fit <- lm(danceability ~ loudness + tempo, data = clean_sp)
summary(lm_fit)
```

Call:

```
lm(formula = danceability ~ loudness + tempo, data = clean_sp)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.4618	-0.1092	0.0078	0.1148	0.3882

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.7546343	0.0477475	15.805	< 2e-16 ***
loudness	-0.0104633	0.0023392	-4.473	9.41e-06 ***
tempo	-0.0013111	0.0002879	-4.554	6.53e-06 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.153 on 537 degrees of freedom


Multiple R-squared: 0.08336, Adjusted R-squared: 0.07994
F-statistic: 24.42 on 2 and 537 DF, p-value: 7.089e-11

Diagnostics


```
# Residuals vs Fitted  
plot(lm_fit, which = 1)
```




```
# Q-Q plot of residuals  
plot(lm_fit, which = 2)
```



```
# Scale-Location
plot(lm_fit, which = 3)
```



```
# Residuals vs Leverage
plot(lm_fit, which = 5)
```

```
# Variance Inflation Factor (multicollinearity)
vif(lm_fit)
```

```
loudness    tempo
1.02822    1.02822
```

```
# Independence of residuals
durbinWatsonTest(lm_fit)
```

```
lag Autocorrelation D-W Statistic p-value
1      0.3796188      1.239978      0
Alternative hypothesis: rho != 0
```

```
# Normality test on residuals
shapiro.test(residuals(lm_fit))
```

Shapiro-Wilk normality test

```
data: residuals(lm_fit)
W = 0.99366, p-value = 0.02307
```

Chi-Square Test of Independence (genre vs danceability)

```
chi_sp <- filter(clean_sp, genre != "Pop")
dance_cat <- ifelse(chi_sp$danceability > median(chi_sp$danceability),
  "Danceable", "Not Danceable")

# Observed counts
obs_tbl <- table(dance_cat, chi_sp$genre)
obs_tbl
```

dance_cat	Dark	Trap	dnb	Emo	hardstyle	Hiphop	psytrance	Rap	RnB	techhouse	
Danceable		20	7	1		2	20	21	15	7	40
Not Danceable		23	31	19		46	18	29	2	13	4


dance_cat	techno	trance	trap	Trap	Meta	Undergrou
Danceable	27	9	11		20	69
Not Danceable	5	31	25		13	10

```
# Expected counts
chi_res <- chisq.test(obs_tbl)
chi_res$expected
```


dance_cat	Dark	Trap	dnb	Emo	hardstyle	Hiphop	psytrance	Rap	RnB	techhouse	
Danceable		21.5	19	10		24	19	25	8.5	10	22
Not Danceable		21.5	19	10		24	19	25	8.5	10	22

dance_cat	techno	trance	trap	Trap	Meta	Undergrou
Danceable	16	20	18		16.5	39.5
Not Danceable	16	20	18		16.5	39.5

```
# Genre bar charts (unfiltered vs filtered)
ggplot(clean_sp, aes(x = genre)) +
  geom_bar(fill = "red") +
  ggtitle("Genre Bar Chart (Unfiltered)") +
  xlab("Genre") + ylab("Count")
```



```
ggplot(chi_sp, aes(x = genre)) +
  geom_bar(fill = "green") +
  ggtitle("Genre Bar Chart (Filtered)") +
  xlab("Genre") + ylab("Count")
```



```
# Chi-square test  
chi_res
```

Pearson's Chi-squared test

data: obs_tbl
X-squared = 192.7, df = 13, p-value < 2.2e-16